

# HW2 Report

## Part 1: Infrastructure Setup

### Code Implementation

- **`train.py`**: Implements training loop with:
  - Character-level tokenization
  - Loss masking (only compute loss after '=' token)
  - Support for masking first N tokens ( `-mask_first_n` flag)
  - Automatic saving of final model and training curves
- **`inference.py`**: Implements text generation with:
  - Model loading from checkpoints
  - Temperature and top-k sampling
  - Interactive mode for testing

### Modifications from Original Codebase

1. **Custom Dataset Class**: `AlgorithmicDataset` with loss masking for equation format
2. **Character Tokenizer**: Simple tokenizer for mathematical expressions
3. **Accuracy Metrics**: Added accuracy computation alongside loss
4. **Visualization**: Automatic generation of training curves


### Sanity Check Results

```
# Test 1: Basic memorization
python train.py --data_dir data/sanity_check --out_dir out/sanity_check --n_layer 1 --n_embd 32 --n_head 4 --max_steps 1000 --log_interval 100 --batch_size 64 --learning_rate 3e-4 --seed 42 --eval_points_per_decade 16


# Test 2: Masked first 3 tokens
python train.py --data_dir data/sanity_check --out_dir out/sanity_check_masked --n_layer 1 --n_embd 32 --n_head 4 --max_steps 1000 --log_interval 100 --batch_size 64 --learning_rate 3e-4 --seed 42 --eval_points_per_decade 16 --mask_first_n 3
```

Both tests successfully converged with train loss  $\rightarrow 0$ . You can run the `inference.py` on the model checkpoint for two sanity checks.

## Sanity Curves



## Masked Sanity Curves



## Challenges Faced

- Debugging loss masking

## Part 2.1: Data Generation

### Process Description

We generate complete datasets for modular arithmetic:

- ```
# For each operation and modulus p:  
- Addition:  $a + b = c$  where  $c = (a + b) \% p$   
- Subtraction:  $a - b = c$  where  $c = (a - b) \% p$   
- Division:  $a / b = c$  where  $a = (b * c) \% p$ 
```

Division uses the reformulation where  $a/b \equiv c \pmod{p}$  means  $a \equiv b \times c \pmod{p}$ .

## Dataset Statistics

| Task        | Modulus | Total Examples | Train | Val   | Test  |
|-------------|---------|----------------|-------|-------|-------|
| Addition    | 97      | 9,409          | 6,586 | 1,411 | 1,412 |
| Subtraction | 97      | 9,409          | 6,586 | 1,411 | 1,412 |
| Division    | 97      | 9,312          | 6,518 | 1,396 | 1,398 |
| Addition    | 113     | 12,769         | 8,938 | 1,915 | 1,916 |
| Subtraction | 113     | 12,769         | 8,938 | 1,915 | 1,916 |

## Part 2.2: Addition and Subtraction Experiments ✓

### Configuration

```
{
  'max_steps': 100000,
  'learning_rate': 1e-3,
  'log_interval': 1000,
  'n_embd': 128,
  'n_head': 4,
  'eval_points_per_decade': 16,
  'batch_size': 64
}
```



## Results (3 Random Seeds: 42, 123, 456)

### 1-Layer Model Results



| Task      | Metric           | Seed 42 | Seed 123 | Seed 456 |
|-----------|------------------|---------|----------|----------|
| Add p=97  | Final Train Loss | 0%      | 0%       | 0%       |
| Add p=97  | Final Train Acc  | 100%    | 100%     | 100%     |
| Add p=97  | Final Test Acc   | 99.04%  | 99.36%   | 99.07%   |
| Add p=113 | Final Train Loss | 0%      | 0%       | 0%       |
| Add p=113 | Final Train Acc  | 100%    | 100%     | 100%     |
| Add p=113 | Final Test Acc   | 97.27%  | 97.26%   | 97.78%   |
| Sub p=97  | Final Train Loss | 0%      | 0%       | 0%       |
| Sub p=97  | Final Train Acc  | 100%    | 100%     | 100%     |
| Sub p=97  | Final Test Acc   | 97.08%  | 91.92%   | 97.15%   |
| Sub p=113 | Final Train Loss | 2%      | 0%       | 0%       |
| Sub p=113 | Final Train Acc  | 100%    | 100%     | 100%     |

|           |                |        |        |        |
|-----------|----------------|--------|--------|--------|
| Sub p=113 | Final Test Acc | 91.49% | 94.06% | 93.54% |
|-----------|----------------|--------|--------|--------|



Add mod 97, 1 layer seed 42





Add mod 113, 1 layer seed 42



Sub mod 97, 1 layer seed 42





Sub mod 113, 1 layer seed 42





## 2-Layer Model Results

| Task      | Metric           | Seed 42 | Seed 123 | Seed 456 |
|-----------|------------------|---------|----------|----------|
| Add p=97  | Final Train Loss | 0%      | 0%       | 0%       |
| Add p=97  | Final Train Acc  | 100.0%  | 100.0%   | 100.0%   |
| Add p=97  | Final Test Acc   | 99.78%  | 99.96%   | 99.78%   |
| Add p=113 | Final Train Loss | 0%      | 0%       | 0%       |
| Add p=113 | Final Train Acc  | 100%    | 100%     | 100%     |
| Add p=113 | Final Test Acc   | 98.49%  | 98.04%   | 98.34%   |
| Sub p=97  | Final Train Loss | 0%      | 0%       | 0%       |
| Sub p=97  | Final Train Acc  | 100%    | 100%     | 100%     |
| Sub p=97  | Final Test Acc   | 99.2%   | 99.31%   | 99.64%   |
| Sub p=113 | Final Train Loss | 0%      | 0%       | 1.8%     |
| Sub p=113 | Final Train Acc  | 100%    | 100%     | 100%     |
| Sub p=113 | Final Test Acc   | 98.46%  | 98.44%   | 97.87%   |



Add mod 97, 2 layer seed 42





Add mod 113, 2 layer seed 42



Sub mod 97, 2 layer seed 42



Sub mod 113, 2 layer seed 42




## Model Checkpoints

Final models saved at: `out/[task]_mod[p]_layer[n]_seed[n]/final_model.pt`

## Part 2.3: Grokking

We trained on the modulo division task for  $p = 97$ , and the training curves are as follows:



### Detailed training configurations:

```
{
  'n_layer': 2,
  'n_embd': 128,
  'n_head': 4,
  'batch_size': 512,
  'learning_rate': 1e-3,
  'weight_decay': 1.0,
  'beta_1': 0.9,
  'beta_2': 0.98
  'max_steps': 100000
}
```

### Grokking Results

- Final **validation loss**: 0.0972, **validation accuracy**: 0.9913.
- Final **training loss**: 0.0160, **training accuracy**: 1.0000

### Model Checkpoint

`out/divide_mod97_layer2_batch512/final_model.pt`

### Inference Instructions

```
python inference.py --checkpoint out/divide_mod97_layer2_batch512/final_model.pt --prompts "30/74=" --max_new_tokens 2
```

## Part 2.4: Ablation Study

### Motivation

We investigated how batch size affects grokking speed and reliability, as batch size influences gradient noise and learning dynamics.

### Experimental Setup

- **Fixed:**

```
{
  'n_layer': 2,
  'n_embd': 128,
  'n_head': 4,
  'learning_rate': 1e-3,
  'weight_decay': 1.0,
  'beta_1': 0.9,
  'beta_2': 0.98
  'max_steps': 100000
  'p': 97
}
```

- **Varied:** batch\_size  $\in \{64, 128, 256, 512\}$
- We tested across three different seeds to ensure reliability. For simplicity we report only the seed=42 plots here, but plots for other seeds are available at `out/divide_mod97_layer2_<suffix>` . Result path to batch size mapping is as follows:



| Batch Size | Path                                                         |
|------------|--------------------------------------------------------------|
| 64         | <code>out/divide_mod97_layer2_seed{seed}_batch{batch}</code> |
| 128        | <code>out/divide_mod97_layer2_seed{seed}_batch{batch}</code> |
| 256        | <code>out/divide_mod97_layer2_seed{seed}_batch{batch}</code> |
| 512        | <code>out/divide_mod97_layer2_seed{seed}_batch{batch}</code> |

### Results



| Batch Size | Steps to Grok* | Final Training Loss | Final Training Accuracy | Final Test Accuracy |
|------------|----------------|---------------------|-------------------------|---------------------|
| 64         | Not Grokking** | 1.7755              | 18.25%                  | 17.06%              |
| 128        | ~9000          | 0.0191              | 100%                    | 97.85%              |
| 256        | ~10000         | 0.0032              | 100%                    | 99.24%              |
| 512        | ~800           | 0.0007              | 100%                    | 99.57%              |





Batch size=64





Batch size=128



Batch size=256



Batch size=512



\*We mark the steps-to-grok as the number of steps the model takes to reach a generally high and stable test accuracy.

\*\*We find that on batch size 64 there's no sign of grokking at all, i.e. the test accuracy keeps bouncing in a very low range and do not present signs of increasing.